

Liangyu Zhao

Email: liangyu@cs.washington.edu

Website: <https://liangyuzhao.me/>

Research Interests Machine learning systems, distributed systems, collective communications; broadly speaking, I am interested in applying mathematical techniques to design and build efficient, scalable computer systems.

Education **University of Washington** Seattle, WA
Ph.D. in Computer Science 2021 – Present

Direction: Systems & Networking
Advisor: Prof. Arvind Krishnamurthy

University of Washington Seattle, WA
M.S. in Computer Science (incomplete) 2020 – 2021

University of Washington Seattle, WA
B.S. in Computer Science,
B.S. in Applied & Computational Mathematical Sciences
(Discrete Math and Algorithms track) 2015 – 2020

Industry Experience **Meta**, Superintelligence Lab (MSL) Infra Menlo Park, CA
Research Scientist Intern Jun – Oct 2025
Mentors: Bingzhe Liu, Liang Luo, Amar Phanishayee
Improved recommendation model QPS via novel embedding table sharding and pipeline parallelism.

NVIDIA, Applied Deep Learning Research (ADLR) Redmond, WA
Research Intern Mar – Jun 2025
Mentors: Vijay Anand Korthikanti & Deepak Narayanan
Designed and optimized customized communication kernels for Megatron-LM.

Microsoft Research, Research in Software Engineering (RiSE) Redmond, WA
Part-Time Researcher Jul – Nov 2024

Microsoft Research, Research in Software Engineering (RiSE) Redmond, WA
Research Intern Jun – Sep 2023
Mentor: Saeed Maleki
Developed collective communication algorithms for multi-node GPU clusters.

ByteDance, AI-Lab Bellevue, WA
Research Intern, ML System Jul – Oct 2020
Mentor: Yibo Zhu
Worked on automatic learning-rate schedule.

Microsoft , Azure Compute Core Software Engineer Intern	Redmond, WA Sep – Dec 2019
Google , Ads Infra Software Engineer Intern	Mountain View, CA Jun – Sep 2019
Microsoft , Azure Compute Core Software Engineer Intern	Redmond, WA Jun – Aug 2018
Zap Surgical Systems Software Engineer Intern	San Carlos, CA Jun – Sep 2017

Publications

ForestColl: Throughput-Optimal Collective Communications on Heterogeneous Network Fabrics

Liangyu Zhao, Saeed Maleki, Yuanhong Wang, Zezhou Wang, Ziyue Yang,
Hossein Pourreza, Arvind Krishnamurthy
arXiv preprint, in submission

FLASH: Fast All-to-All Communication in GPU Clusters
Yiran Lei, Dongjoo Lee, **Liangyu Zhao**, Daniar Kurniawan, Chanmyeong Kim,
Heetaek Jeong, Changsu Kim, Hyeonseong Choi, Liangcheng Yu, Arvind Krishnamurthy, Justine Sherry, Eriko Nurvitadhi
arXiv preprint, in submission

Tactic: Adaptive Sparse Attention with Clustering and Distribution Fitting for Long-Context LLMs
Kan Zhu*, Tian Tang*, Qinyu Xu*, Yile Gu, Zhichen Zeng, Rohan Kadekodi,
Liangyu Zhao, Ang Li, Arvind Krishnamurthy, Baris Kasikci
arXiv preprint, in submission

NanoFlow: Towards Optimal Large Language Model Serving Throughput
Kan Zhu, Yufei Gao, Yilong Zhao, **Liangyu Zhao**, Gefei Zuo, Yile Gu, Dedong Xie, Tian Tang, Qinyu Xu, Zihao Ye, Keisuke Kamahori, Chien-Yu Lin, Ziren Wang, Stephanie Wang, Arvind Krishnamurthy, Baris Kasikci
USENIX Symposium on Operating Systems Design and Implementation (OSDI '25)

Efficient Direct-Connect Topologies for Collective Communications
Liangyu Zhao, Siddharth Pal, Tapan Chugh, Weiyang Wang, Jason Fantl,
Prithwish Basu, Joud Khoury, Arvind Krishnamurthy
USENIX Symposium on Networked Systems Design and Implementation (NSDI '25)

Rethinking Machine Learning Collective Communication as a Multi-Commodity Flow Problem

Xuting Liu, Behnaz Arzani, Siva Kesava Reddy Kakarla, **Liangyu Zhao**, Vincent Liu, Miguel Castro, Srikanth Kandula, Luke Marshall

ACM Special Interest Group on Data Communication (SIGCOMM '24)

Efficient all-to-all Collective Communication Schedules for Direct-connect Topologies

Prithwish Basu, **Liangyu Zhao**, Jason Fantl, Siddharth Pal, Arvind Krishnamurthy, Joud Khoury

International Symposium on High-Performance Parallel and Distributed Computing (HPDC '24)

AutoLRS: Automatic Learning-Rate Schedule by Bayesian Optimization on the Fly

Yuchen Jin, Tianyi Zhou, **Liangyu Zhao**, Yibo Zhu, Chuanxiong Guo, Marco Canini, Arvind Krishnamurthy

International Conference on Learning Representations (ICLR '21)

Nexus: A GPU Cluster Engine for Accelerating DNN-Based Video Analysis

Haichen Shen, Lequn Chen, Yuchen Jin, **Liangyu Zhao**, Bingyu Kong, Matthai Philipose, Arvind Krishnamurthy, Ravi Sundaram

ACM Symposium on Operating Systems Principles (SOSP '19)

Invited Talks

Efficient Direct-Connect Topologies for Collective Communications

➤ USENIX NSDI '25 April, 2025

➤ ACE Liaison Meeting Theme 3

ACE Center for Evolvable Computing January, 2025

➤ Future of Cloud Infrastructure (FOCI) Annual Symposium

University of Washington October, 2023

➤ Harvard Cloud Networking and Systems Group

Harvard University July, 2023

ForestColl: Throughput-Optimal Collective Communications on Heterogeneous Network Fabrics

➤ Network and Mobile System Group

Massachusetts Institute of Technology July, 2025

➤ Distributed Systems Laboratory (DSL) Seminar

University of Pennsylvania November, 2024

➤ NLP Reading Group

NVIDIA November, 2024

➤ Paul G. Allen School Annual Research Showcase

University of Washington October, 2024

➤ Research in Software Engineering (RiSE)

Microsoft Research August, 2024

➤ ByteDance

August, 2024

➤ AMD Research

July, 2024