

# ForestColl: Efficient Collective Communications on Heterogeneous Network Fabrics

- Originally a topic in HPC, it is now extensively used for gradient, parameter, and activation synchronization in distributed ML training and inferencing.
- Allgather** is an operation where every node/GPU broadcasts a shard of data.
  - reduce-scatter = *reversed* allgather
  - allreduce = reduce-scatter + allgather

Before			After		
Reduce-Scatter					
Node 0	Node 1	Node 2	Node 0	Node 1	Node 2
$S_0^{(0)}$	$S_0^{(1)}$	$S_0^{(2)}$	$\bigoplus_i S_0^{(i)}$		
$S_1^{(0)}$	$S_1^{(1)}$	$S_1^{(2)}$		$\bigoplus_i S_1^{(i)}$	
$S_2^{(0)}$	$S_2^{(1)}$	$S_2^{(2)}$			$\bigoplus_i S_2^{(i)}$
Allgather					
Node 0	Node 1	Node 2	Node 0	Node 1	Node 2
$S_0^{(0)}$			$S_0^{(0)}$	$S_0^{(0)}$	$S_0^{(0)}$
	$S_1^{(1)}$		$S_1^{(1)}$	$S_1^{(1)}$	$S_1^{(1)}$
		$S_2^{(2)}$	$S_2^{(2)}$	$S_2^{(2)}$	$S_2^{(2)}$
Allreduce					
Node 0	Node 1	Node 2	Node 0	Node 1	Node 2
$S_0^{(0)}$	$S_0^{(1)}$	$S_0^{(2)}$	$\bigoplus_i S_0^{(i)}$	$\bigoplus_i S_0^{(i)}$	$\bigoplus_i S_0^{(i)}$
$S_1^{(0)}$	$S_1^{(1)}$	$S_1^{(2)}$	$\bigoplus_i S_1^{(i)}$	$\bigoplus_i S_1^{(i)}$	$\bigoplus_i S_1^{(i)}$
$S_2^{(0)}$	$S_2^{(1)}$	$S_2^{(2)}$	$\bigoplus_i S_2^{(i)}$	$\bigoplus_i S_2^{(i)}$	$\bigoplus_i S_2^{(i)}$

We aim to derive efficient communication schedules for any given network topology.

- **Heterogeneity:** today's ML network architectures are highly diverse and heterogeneous.
- **Scalability:** optimizing aggregation and multicast traffic requires strict data dependency, often results in NP-hard discrete optimization.

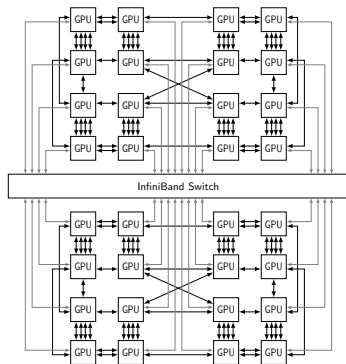


Figure: AMD MI250 Topology

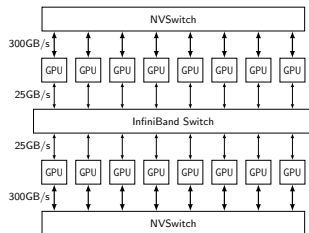


Figure: NVIDIA DGX A100 Topology

# of nodes	4	9	16	25	36
SCCL [PPoPP '21]	0.61s	1.00s	60s	3286s	$>10^4$ s
TACCL [NSDI '23]	0.45s	67.8s	1801s	1802s	n/a

Table: Generation Time on 2D Torus ( $n \times n$ )

ForestColl: construct a set of spanning trees with  $k$  trees rooted at each node/GPU.

- Each tree broadcasts  $1/k$  of the data from each root.
- **Performance:** the trees achieve mathematically **minimum overlap/congestion**.
- **Scalability:** computation is in **strongly polynomial time**.

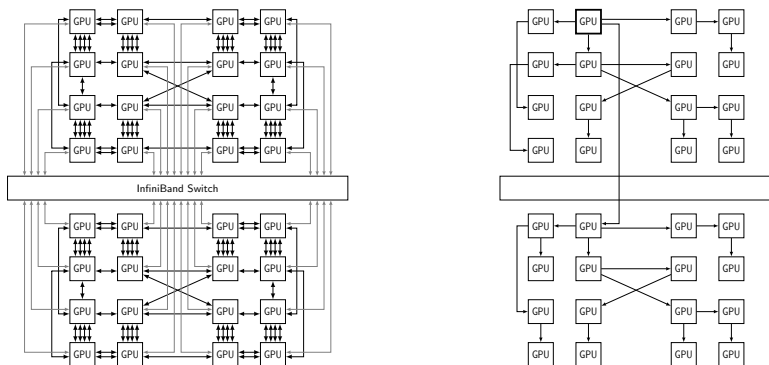


Figure: 2-Node AMD MI250

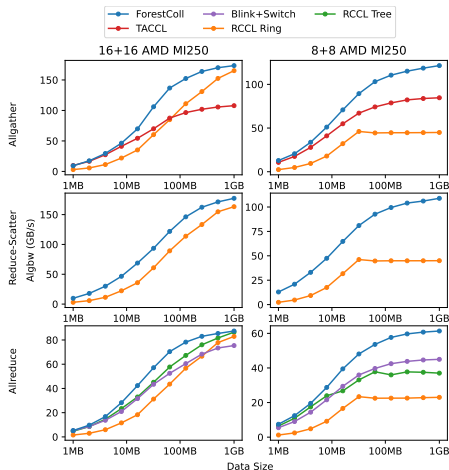
	SCCL	TACCL	BFB	Blink	TE-CCL	<b>ForestColl</b>
Switch-based Network	×	✓	×	×	✓	✓
Optimal Schedule	✓	×	×	×	×	✓
Scalable Runtime	×	×	✓	✓	×	✓

**Table:** Comparison of schedule generation methods.

Compared to previous work:

- ForestColl does not require input parameters such as the number of chunks or chunk size, eliminating the need for parameter sweeps.
- While converting switch topologies into equivalent switchless logical topologies, ForestColl ensures no compromise to performance.
- Blink constructs trees rooted at a single node. Thus, Blink does not support allgather/reduce-scatter and performs allreduce as reduce+broadcast, which can suffer from a bottleneck at a single GPU.

# Evaluation



16+16 Setting:

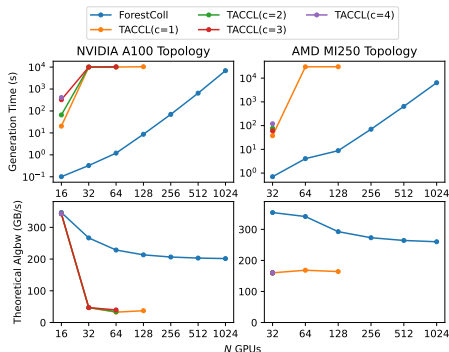
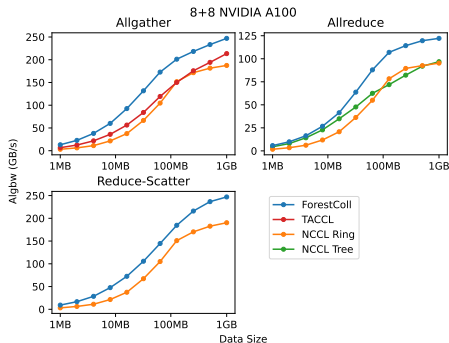
- ForestColl outperforms RCCL by 37%, 32%, and 13% on average in allgather, reduce-scatter, and allreduce.

8+8 Setting (half of the GPUs per node):

- ForestColl outperforms RCCL by 2.52x, 2.29x, and 1.47x on average in allgather, reduce-scatter, and allreduce.

## ForestColl is

- on average, 52%, 38%, and 30% faster than NCCL on 8+8 A100 in allgather, reduce-scatter, and allreduce, respectively.
- orders of magnitude faster schedule generation than TACCL.



**Q:** Why not just use rings?

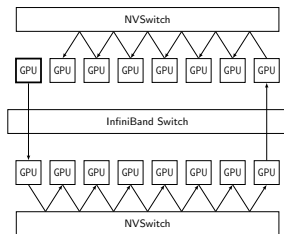


Figure: NCCL Ring

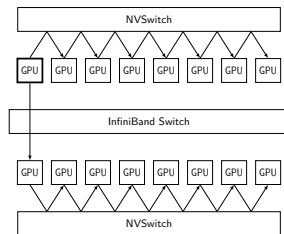


Figure: ForestColl



**Q:** Why not just use rings?

- Intra-node bandwidth is much more abundant than inter-node bandwidth.

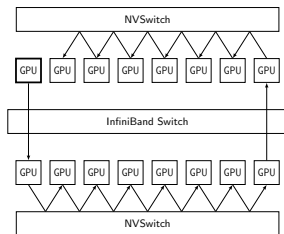


Figure: NCCL Ring

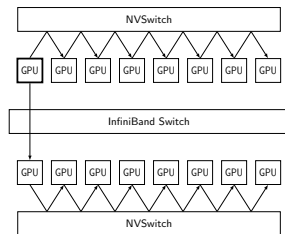


Figure: ForestColl

## Q: Why not just use rings?

- Intra-node bandwidth is much more abundant than inter-node bandwidth.
- Rings often overuse inter-node bandwidth, even though the traffic can be delivered within the node.
  - In ring allgather, every GPU performs a ring broadcast. In total, ring allgather has 2x amount of inter-node traffic as ForestColl.

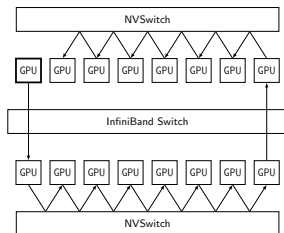


Figure: NCCL Ring

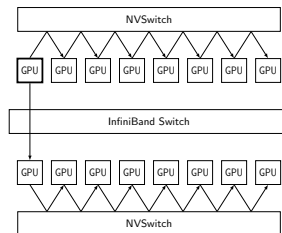


Figure: ForestColl

## Q: Why not just use rings?

- Intra-node bandwidth is much more abundant than inter-node bandwidth.
- Rings often overuse inter-node bandwidth, even though the traffic can be delivered within the node.
  - In ring allgather, every GPU performs a ring broadcast. In total, ring allgather has 2x amount of inter-node traffic as ForestColl.

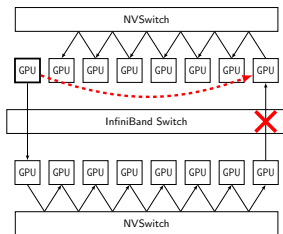


Figure: NCCL Ring

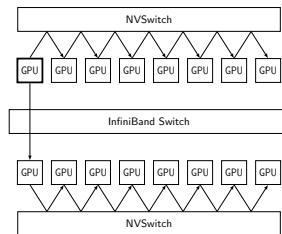


Figure: ForestColl

## Q: Why not just use rings?

- Intra-node bandwidth is much more abundant than inter-node bandwidth.
- Rings often overuse inter-node bandwidth, even though the traffic can be delivered within the node.
  - In ring allgather, every GPU performs a ring broadcast. In total, ring allgather has 2x amount of inter-node traffic as ForestColl.
- **Full Bandwidth Utilization  $\neq$  Optimal Performance**

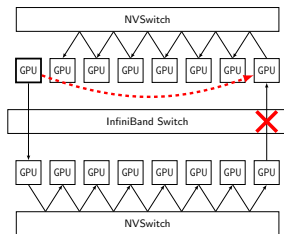


Figure: NCCL Ring

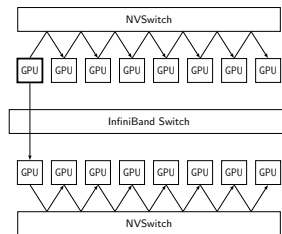
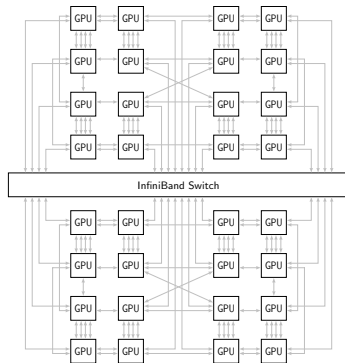


Figure: ForestColl

**Q:** What is the optimal allgather performance given a topology?



**Figure:** 2-Node AMD MI250

**Q:** What is the optimal allgather performance given a topology?

- Previous works often use  $\frac{M}{B} \cdot \frac{N-1}{N}$ .
  - The minimum amount of data needs to be received by a GPU is  $\frac{M}{N}(N-1)$ .
  - The total ingress bandwidth of the GPU is  $B$ .

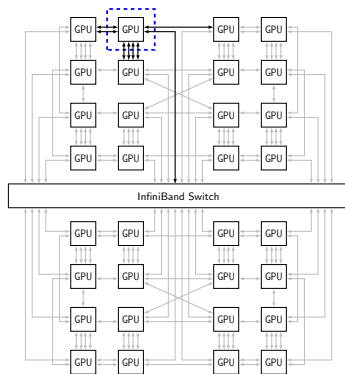


Figure: 2-Node AMD MI250

**Q:** What is the optimal allgather performance given a topology?

- Previous works often use  $\frac{M}{B} \cdot \frac{N-1}{N}$ .
  - The minimum amount of data needs to be received by a GPU is  $\frac{M}{N}(N-1)$ .
  - The total ingress bandwidth of the GPU is  $B$ .
- What if the performance is bounded by some network cut elsewhere?

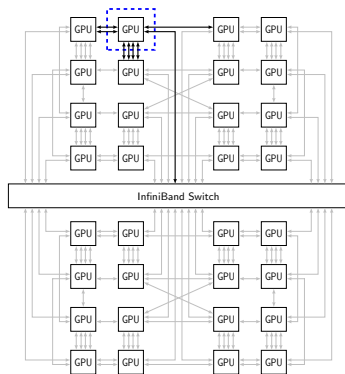


Figure: 2-Node AMD MI250

**Q:** What is the optimal allgather performance given a topology?

- Consider a cut  $S$  from the network:

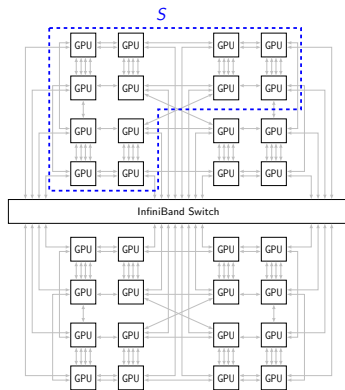


Figure: 2-Node AMD MI250



**Q:** What is the optimal allgather performance given a topology?

- Consider a cut  $S$  from the network:
  - The amount of data needs to travel out of  $S$  is at least  $\frac{M}{N} \cdot |S \cap V_c|$ . ( $V_c$  = set of GPUs)

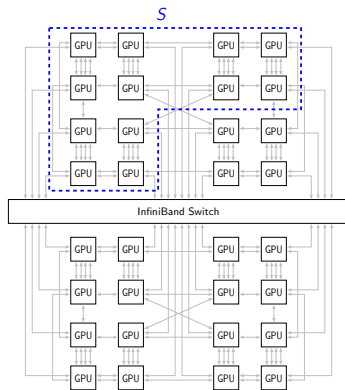


Figure: 2-Node AMD MI250

**Q:** What is the optimal allgather performance given a topology?

- Consider a cut  $S$  from the network:
  - The amount of data needs to travel out of  $S$  is at least  $\frac{M}{N} \cdot |S \cap V_c|$ . ( $V_c$  = set of GPUs)
  - The exiting bandwidth of  $S$  is  $B^+(S)$ .

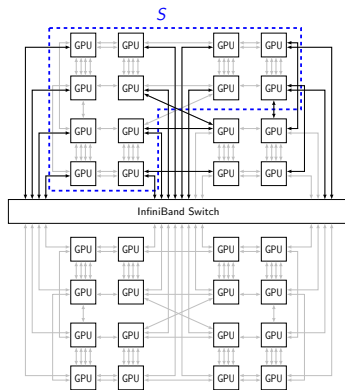


Figure: 2-Node AMD MI250

**Q:** What is the optimal allgather performance given a topology?

- Consider a cut  $S$  from the network:
  - The amount of data needs to travel out of  $S$  is at least  $\frac{M}{N} \cdot |S \cap V_c|$ . ( $V_c$  = set of GPUs)
  - The exiting bandwidth of  $S$  is  $B^+(S)$ .

The allgather runtime is at least  $\frac{M}{N} \cdot \frac{|S \cap V_c|}{B^+(S)}$ .

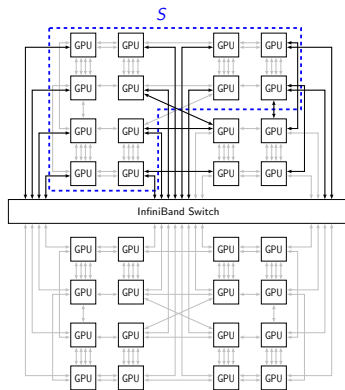


Figure: 2-Node AMD MI250

**Q:** What is the optimal allgather performance given a topology?

- Consider a cut  $S$  from the network:
  - The amount of data needs to travel out of  $S$  is at least  $\frac{M}{N} \cdot |S \cap V_c|$ . ( $V_c$  = set of GPUs)
  - The exiting bandwidth of  $S$  is  $B^+(S)$ .

The allgather runtime is at least  $\frac{M}{N} \cdot \frac{|S \cap V_c|}{B^+(S)}$ .

- The optimal performance is determined by a bottleneck cut  $S^*$  such that

$$\frac{|S^* \cap V_c|}{B^+(S^*)} = \max_{S \subset V, S \not\subseteq V_c} \frac{|S \cap V_c|}{B^+(S)}.$$

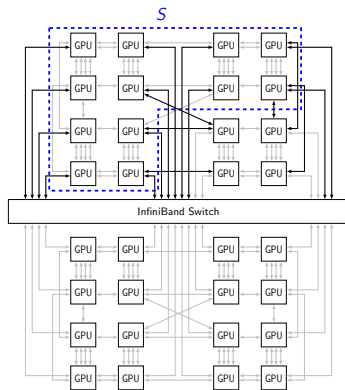


Figure: 2-Node AMD MI250

**Q:** What is the optimal allgather performance given a topology?

- Consider a cut  $S$  from the network:
  - The amount of data needs to travel out of  $S$  is at least  $\frac{M}{N} \cdot |S \cap V_c|$ . ( $V_c$  = set of GPUs)
  - The exiting bandwidth of  $S$  is  $B^+(S)$ .

The allgather runtime is at least  $\frac{M}{N} \cdot \frac{|S \cap V_c|}{B^+(S)}$ .

- The optimal performance is determined by a bottleneck cut  $S^*$  such that

$$\frac{|S^* \cap V_c|}{B^+(S^*)} = \max_{S \subset V, S \not\supseteq V_c} \frac{|S \cap V_c|}{B^+(S)}.$$

ForestColl can generate spanning trees to achieve the above optimality in strongly polynomial time.

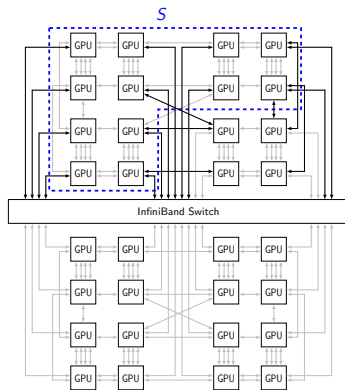


Figure: 2-Node AMD MI250

## AMD MI250:

- When number of nodes  $< 4$ , the ingress bandwidth of an OAM is the bottleneck.
- When number of nodes  $\geq 4$ , the ingress bandwidth of a node is the bottleneck.

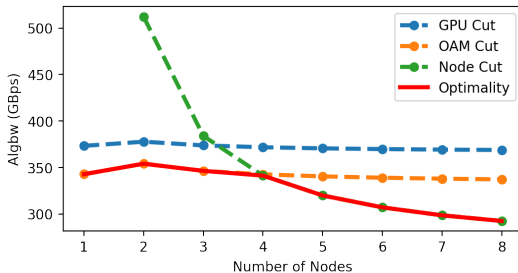
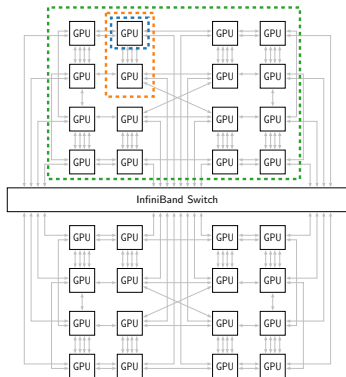
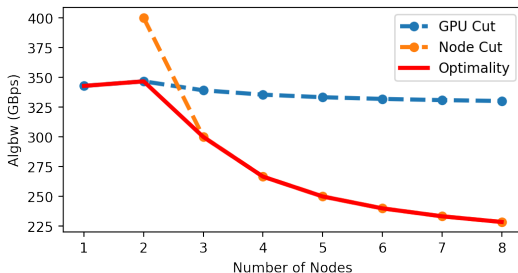
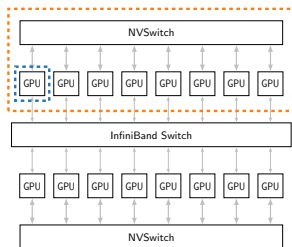


Figure: Optimality and performance bounds from different cuts of AMD MI250 topologies

## NVIDIA DGX A100:

- When number of nodes  $< 3$ , the ingress bandwidth of a GPU is the bottleneck.
- When number of nodes  $\geq 3$ , the ingress bandwidth of a node is the bottleneck.



**Figure:** Optimality and performance bounds from different cuts of NVIDIA DGX A100 topologies

ForestColl is a schedule generation algorithm for collective communications that

- provides **provably optimal** schedule;
- works on **any network topology** (direct-connect or switch topology);
- runs in **strongly polynomial time** (scalable to large number of nodes);
- outperforms state-of-the-art solutions.

Paper: <https://arxiv.org/abs/2402.06787>

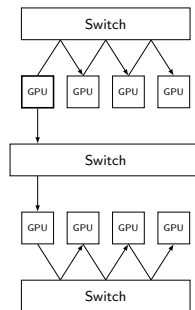
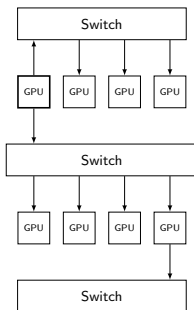
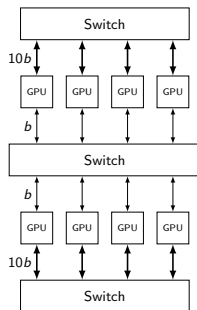
GitHub: <https://github.com/liangyuRain/ForestColl>



# Switch Topology

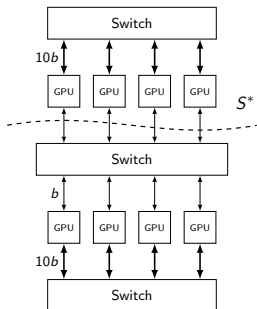
In switch topology, the vertex set consists of **compute nodes** and **switch nodes**.

- **Problem:** allgather is no longer defined by spanning out-trees.
  - Non-Spanning: unnecessary to broadcast data to every switch node.
  - Non-Tree: switch may not be able to multicast.
- **Solution:** convert switch topology into a logical topology without switches.

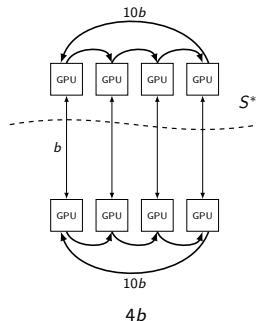
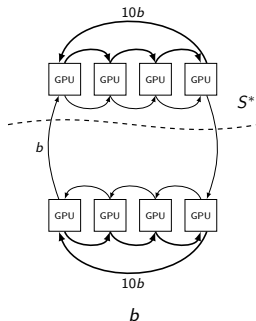


# Edge Splitting

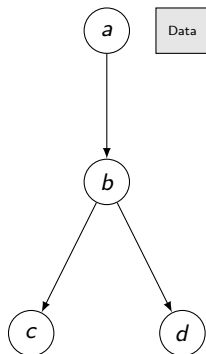
- **Previous work** proposed ways such as unwinding a switch into a ring.
- **Edge Splitting:** for each switch node  $w$ , iteratively choose edges  $(u, w)$ ,  $(w, t)$  and replace them by  $(u, t)$  without sacrificing connectivity.
  - Originally used to prove connectivity properties of Eulerian graph. (Jackson, 1988; Frank, 1988; Bang-Jensen et al., 1995)
  - Now to remove switch nodes without compromising allgather performance.



Cut Bandwidth:  $4b$

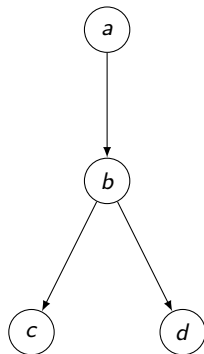


Non-Pipeline Schedule

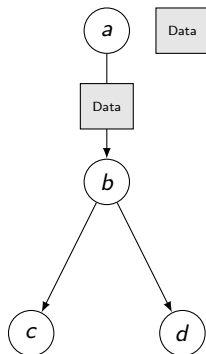


Time Cost: 0

Pipeline Schedule

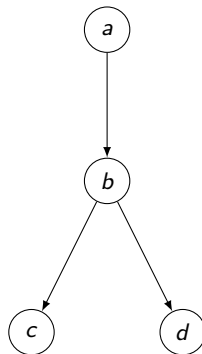


Non-Pipeline Schedule

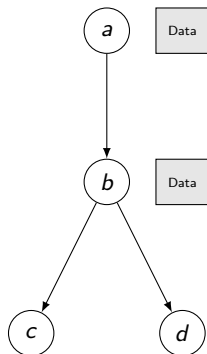


Time Cost: 1

Pipeline Schedule

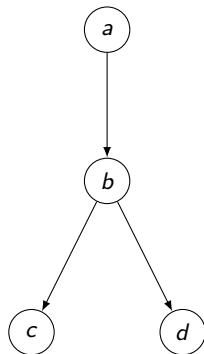


Non-Pipeline Schedule

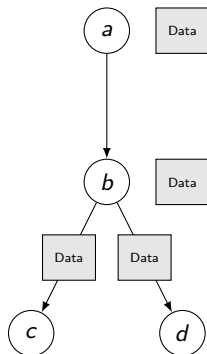


Time Cost: 1

Pipeline Schedule

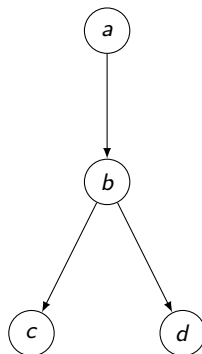


Non-Pipeline Schedule

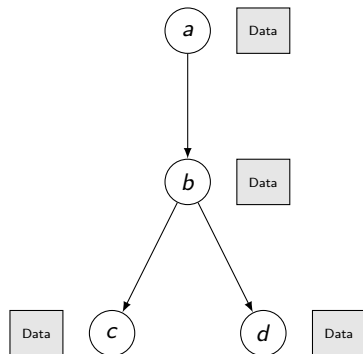


Time Cost: 2

Pipeline Schedule

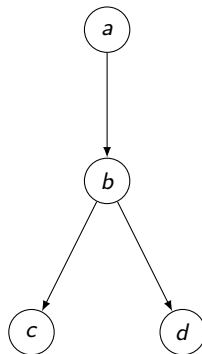


Non-Pipeline Schedule

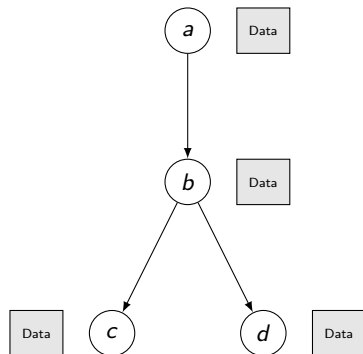


Time Cost: 2

Pipeline Schedule

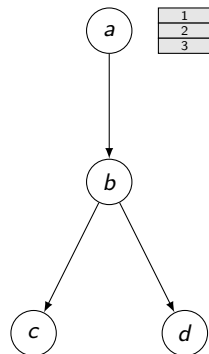


## Non-Pipeline Schedule



Time Cost: 2

## Pipeline Schedule

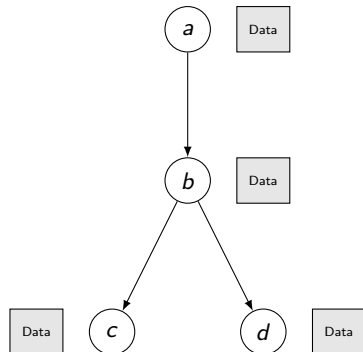


Time Cost: 0



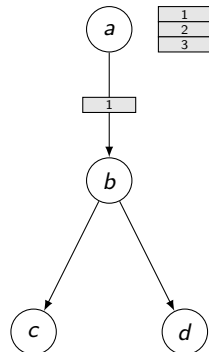
# Pipeline Schedule

Non-Pipeline Schedule



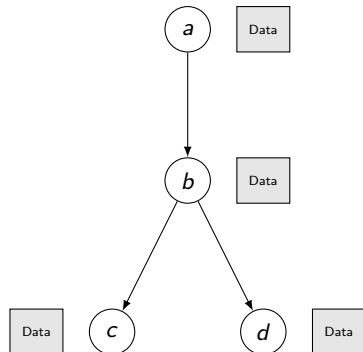
Time Cost: 2

Pipeline Schedule



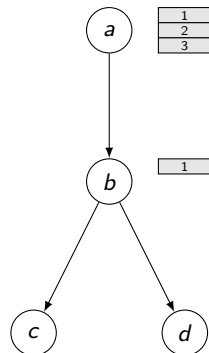
Time Cost:  $1/3$

## Non-Pipeline Schedule



Time Cost: 2

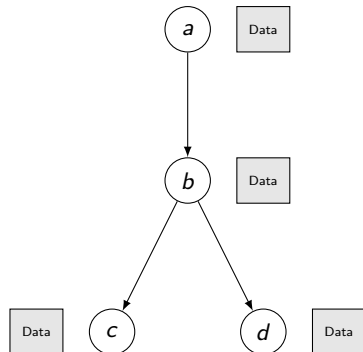
## Pipeline Schedule



Time Cost:  $1/3$

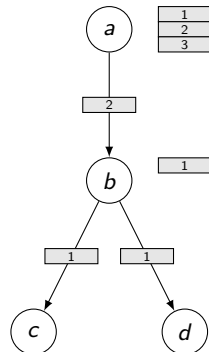
# Pipeline Schedule

Non-Pipeline Schedule



Time Cost: 2

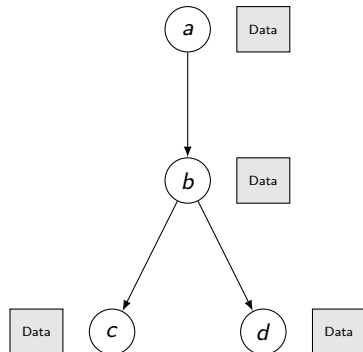
Pipeline Schedule



Time Cost:  $2/3$

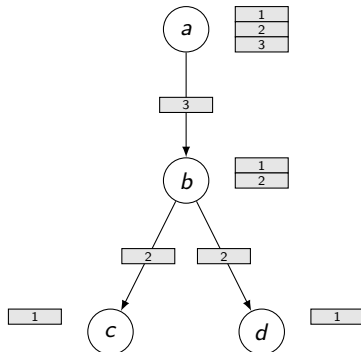
# Pipeline Schedule

Non-Pipeline Schedule



Time Cost: 2

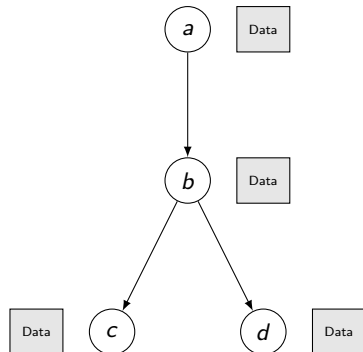
Pipeline Schedule



Time Cost: 3/3

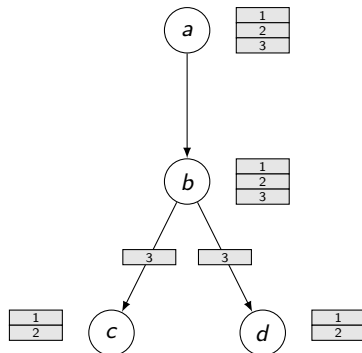
# Pipeline Schedule

Non-Pipeline Schedule



Time Cost: 2

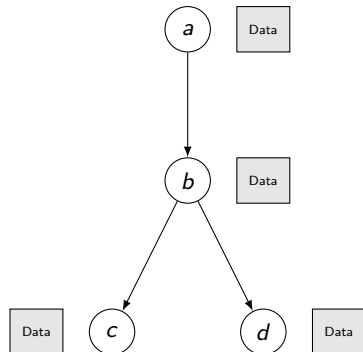
Pipeline Schedule



Time Cost:  $4/3$

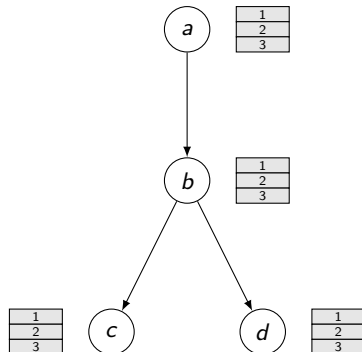
# Pipeline Schedule

Non-Pipeline Schedule



Time Cost: 2

Pipeline Schedule



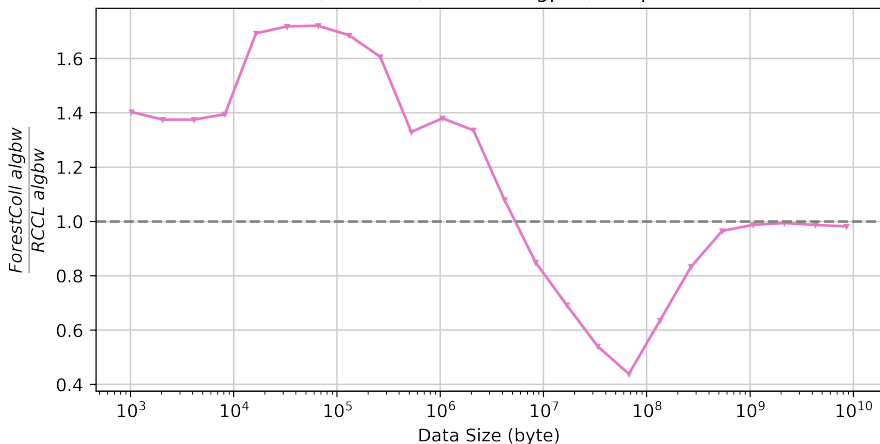
Time Cost:  $4/3$

- ForestColl schedule assumes that data is transmitted as **flows** along the trees rather than through discrete send/rcv steps.
- Ideally, **chunk size** should be as small as possible to enhance bandwidth utilization; however, send/rcv has **overhead** in practice.



# Chunk Size Experiment

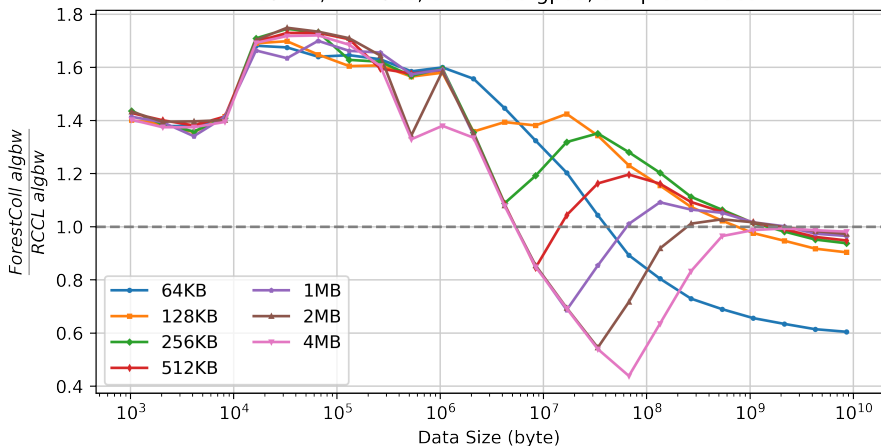
ForestColl Schedule Performance with Default NCCL\_BUFFSIZE  
Allreduce, 2 nodes, 32 MI250 gpus, Simple Protocol





# Chunk Size Experiment

ForestColl Schedule Performance with Different NCCL\_BUFFSIZE  
Allreduce, 2 nodes, 32 MI250 gpus, Simple Protocol



Thank you

Paper: <https://arxiv.org/abs/2402.06787>  
GitHub: <https://github.com/liangyuRain/ForestColl>